

System Context

Generalized Notation Notation (GNN) couples a small, declarative text language for Active Inference generative models to a deterministic processing pipeline that turns each specification into validation, visualization, simulation, and analysis artifacts (Smékal and Friedman 2023). The language gives a model one canonical written form; the pipeline gives that form many executable and graphical realizations. This section describes both halves of the architecture and the way they meet, and it states the mathematical objects the language is obliged to carry.

A GNN model is a plain-text document organized into named sections that together pin down a complete partially observable Markov decision process. The `StateSpaceBlock` declares the variables of the model and their dimensions — hidden states, observations, control factors, and policies — establishing the shape of every tensor that follows. The `Connections` section records the directed and undirected dependencies among those variables, the edges of the underlying factor graph that downstream tools read to lay out diagrams and wire up inference. The full construct vocabulary is catalogued in tbl. 4.

The Generative Model a Specification Denotes I

Every GNN file denotes one object: a discrete state-space generative model over a horizon T , factorized as in eq. 1 following the standard formulation (Da Costa et al. 2020; Smith et al. 2022).

$$P(o_{1:T}, s_{1:T}, \pi) = P(\pi) P(s_1) \prod_{t=1}^T P(o_t | s_t) \prod_{t=2}^T P(s_t | s_{t-1}, \pi) \quad (1)$$

The four matrices named in the `StateSpaceBlock` are exactly the factors of eq. 1, and this correspondence is what makes the notation translatable rather than merely descriptive. The A matrix is the likelihood, a column-stochastic map from hidden states to observation outcomes, given in eq. 2.

The Generative Model a Specification Denotes I

$$P(o_t = i \mid s_t = j) = A_{ij}, \quad \sum_i A_{ij} = 1 \quad (2)$$

The B matrix is the controlled transition dynamics: one column-stochastic slice per action, as in eq. 3.

$$P(s_{t+1} = i \mid s_t = j, u_t = k) = B_{ij}^{(k)}, \quad \sum_i B_{ij}^{(k)} = 1 \quad (3)$$

The C vector holds log-preferences over observations, which enter inference as the biased outcome distribution of eq. 4; the D vector is the prior over initial hidden states, eq. 5.

The Generative Model a Specification Denotes I

$$\tilde{P}(o_t = i) = \sigma(C)_i = \frac{\exp C_i}{\sum_j \exp C_j} \quad (4)$$

$$P(s_1 = i) = D_i, \quad \sum_i D_i = 1 \quad (5)$$

Given those factors, perception is approximate Bayesian inference: a variational posterior $Q(s)$ is fitted by minimizing the variational free energy of eq. 6, which decomposes into a complexity term and an accuracy term (Friston 2010).

$$F[Q] = \underbrace{D_{\text{KL}}[Q(s) \| P(s)]}_{\text{complexity}} - \underbrace{\mathbb{E}_{Q(s)}[\ln P(o | s)]}_{\text{accuracy}} \quad (6)$$

The Generative Model a Specification Denotes I

Action is expected-free-energy minimization: each policy π is scored over future steps τ by eq. 7, whose two terms are the information gain a policy is expected to yield and the extent to which its predicted outcomes match $\mathcal{C}$ (Da Costa et al. 2020).

$$G(\pi) = - \underbrace{\mathbb{E}_{Q(o_\tau, s_\tau | \pi)} [\ln Q(s_\tau | o_\tau, \pi) - \ln Q(s_\tau | \pi)]}_{\text{epistemic value}} - \underbrace{\mathbb{E}_{Q(o_\tau | \pi)} [\ln \tilde{P}(o_\tau)]}_{\text{pragmatic value}} \quad (7)$$

The Generative Model a Specification Denotes I

The policy posterior then combines those scores with the habit prior E declared in the specification, as in eq. 8, and the selected action u is sampled from it.

$$Q(\pi) = \sigma(\ln E - G) \tag{8}$$

Recent work continues to refine how expected-free-energy objectives relate to variational inference and to alternative but equivalent formulations, which is why GNN keeps the objects of eq. 1 through eq. 8 explicit in the notation rather than burying them in backend-specific code (Champion et al. 2024; Nuijten, Lukashchuk, et al. 2026; Nuijten, Laar, et al. 2026). A `ModelParameters` section fixes the scalars those objects depend on — factor cardinalities and precision terms — while a `Time` section declares whether the model is static or dynamic and, if dynamic, how the horizon T and its discretization are

The Generative Model a Specification Denotes I

organized. Because every one of these sections is explicit text, a GNN file is at once human-readable, diffable under version control, and unambiguous to a parser — the property that lets the rest of the pipeline operate deterministically.

The Processing Pipeline I

The pipeline is a fixed sequence of 25 numbered steps, 0–24, each a self-contained stage that consumes the artifacts of its predecessors and writes typed outputs for those that follow. Early steps parse and type-check the GNN text and validate it against the language schema; middle steps render visualizations, export the model to executable backends, and run simulations; later steps perform analysis, reporting, and downstream integration. The data dependencies among the steps form the directed acyclic graph shown in fig. 1, which makes the whole flow inspectable: any artifact can be traced back to the step that produced it and forward to every step that depends on it.

The Processing Pipeline I




Figure 1: The GNN processing pipeline as a directed acyclic graph of numbered steps, from parsing and validation through visualization, execution, and analysis.

The Processing Pipeline I

The per-step responsibilities are enumerated in tbl. 1; each row names a step and the transformation it owns within the 0–24 range.

The Processing Pipeline I

Table 1: The pipeline steps, their thin orchestrator modules, and their purposes, read from `src/STEP_INDEX.md`.

Step	Module	Purpose
0	<code>0_template.py</code>	Pipeline template & initialization
1	<code>1_setup.py</code>	Environment setup & UV dependency install
2	<code>2_tests.py</code>	Test suite execution (pytest)
3	<code>3_gnn.py</code>	GNN file discovery & multi-format parsing
4	<code>4_model_registry.py</code>	Model versioning & registry management
5	<code>5_type_checker.py</code>	GNN type validation & resource estimation
6	<code>6_validation.py</code>	Consistency & semantic quality checking
7	<code>7_export.py</code>	Multi-format export (JSON, XML, GraphML, GEXF, Pickle)
8	<code>8_visualization.py</code>	Graph & matrix visualization generation
9	<code>9_advanced_viz.py</code>	Interactive / advanced visualization (Plotly, D3)
10	<code>10_ontology.py</code>	Active Inference ontology processing & validation
11	<code>11_render.py</code>	Code generation for simulation frameworks
12	<code>12_execute.py</code>	Execute rendered simulation scripts
13	<code>13_llm.py</code>	LLM-enhanced analysis & model interpretation
14	<code>14_ml_integration.py</code>	Machine learning integration & model training
15	<code>15_audio.py</code>	Audio sonification generation (SAPF)
16	<code>16_analysis.py</code>	Statistical analysis & cross-simulation aggregation

The Processing Pipeline II

17	17_integration.py	System integration & cross-module coordination
18	18_security.py	Security validation & generated code scanning
19	19_research.py	Research tools & literature references
20	20_website.py	Static HTML website generation
21	21_mcp.py	Model Context Protocol processing & tool registration
22	22_gui.py	Interactive GNN constructor GUI
23	23_report.py	Comprehensive analysis report generation
24	24_intelligent_analysis.py	AI-powered pipeline analysis & executive reports

The Processing Pipeline I

This staged design keeps the architecture modular. The implementation is organized into 31 source packages, one cluster of responsibilities per concern, and is documented across 616 documentation files so that each step's contract, inputs, and outputs are specified independently of the others. New backends or analyses attach to the graph by declaring their dependencies rather than by editing a monolith, and the deterministic step ordering means a model processed today yields the same artifacts when reprocessed tomorrow.

The Triple Play I

The reason for separating a single written language from a multi-stage pipeline is the design goal GNN calls the Triple Play: one model specification, three coordinated modes of existence. The same GNN text is simultaneously a human-readable model description, a set of graphical visualizations of its state space and factor structure, and an executable cognitive model that can be run as a simulation. fig. 2 depicts these three faces and the shared specification at their center.

The Triple Play I

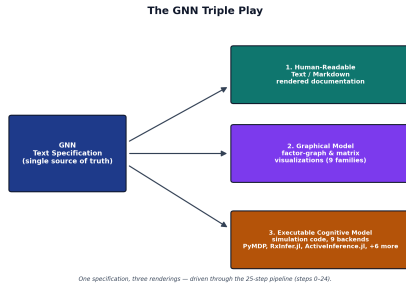


Figure 2: The Triple Play: a single GNN specification rendered as readable text, as graphical visualizations, and as an executable model.

Each face is generated from the same source, so they cannot drift apart. The text is the contract a researcher reads and reviews; the visualizations expose the factor-graph structure for inspection and communication; and the executable rendering lets the identical model be simulated across inference frameworks, connecting the written specification to libraries such as pymdp for discrete Active Inference (Heins et al. 2022) and to the broader tooling ecosystem for compositional model representation (Felice et al. 2021). Because all three derive from one parsed document, GNN turns a model from a static artifact into a live object that can be read, seen, and run without re-encoding it for each purpose (Smith et al. 2022).